

RISK AND RESILIENCE BOOTCAMP





WORKFORCE
DEVELOPMENT



SYSTEMS THINKING

This module introduces some of the basic ideas of systems thinking

- Defining systems thinking
- How it applies to risk analysis
- Cascading failures
- Managing complexity



SYSTEMS THINKING

- Systems thinking is a way of seeing and talking about reality
 - It emphasizes wholes, interrelationships and patterns of change rather than just isolated parts
 - It is a mindset that focuses on how everything is connected
 - Coupled with a set of tools/techniques that help make sense of complexity
 - For example: causal loop diagrams, behaviour-over-time graphs, stock & flow models
 - Doesn't just look at one system or one process in isolation
 - Asks instead
 - How do the people, the process, and the systems interact?
 - What feedback loops, dependencies, delays or unintended effects arise when these interplay?

SYSTEMS THINKING

- Why systems thinking matters
 - IT systems, business processes, and human actors are deeply intertwined
 - A change in one often shifts behaviour in others, sometimes in unexpected ways
 - Without systems thinking
 - There is the risk of designing controls or modifying processes in a silo, missing side-effects in other processes and missing hidden dependencies that might cause cascading failures
 - Systems thinking helps identify emergent behaviours
 - Specifically that the whole is greater than the sum of parts
 - For example, when a process runs across multiple teams and systems, the risk is not just each part failing, it's how their interactions can produce cascading or amplifying effects

SYSTEMS THINKING

- Using systems allows an analyst to
 - Map feedback loops, both positive and negative
 - Recognize delays and time lags between action and effect
 - See non-linearities where a small change can triggers much larger changes
 - Understand dependencies among people, process and systems
 - Move the analysis of causes
 - From "What happened in system X?"
 - To "Why did the interaction across systems/processes/people cause the failure?"

SYSTEMS TRIAD

- people-process-systems triad
 - People
 - Roles, communication, culture, decision-making, skill levels, hand-offs
 - Process
 - Flow of tasks, decision points, hand-offs, controls, workflows, exceptions
 - Systems (IT / technical)
 - Software, hardware, interfaces, data flows, automation, integrations
- Systems thinking
 - Emphasizes how these three overlap and interact.
 - Example:
 - A process may call for a manual hand-off (people), but if the system interface is slow (systems), it may cause delay (process) which leads to human workaround, raising risk (people) of human error

APPLYING SYSTEMS THINKING

- Define the boundary of your system
 - Identify what “system” you are analyzing
 - For example: “vendor on-boarding process + IT systems + stakeholders”
 - Clarify what’s inside, what’s outside, and how it interacts with the environment.
- Map the elements and their relationships
 - Use diagrams (causal loop diagrams, stock/flow diagrams, behaviour-over-time) to visualise how people, processes and systems interact
- Identify feedback loops and delays:
 - For example
 - “System slows → manual work increases → backlog grows → errors increase → system redesign becomes urgent”

APPLYING SYSTEMS THINKING

- Look for unintended consequences
 - When you change one part what are impacts elsewhere
 - For example:
 - If we automate a manual approval, what happens elsewhere?
 - Are there other processes that are affected downstream? Does this impact regulatory compliance in other processes?
- Analyze the dynamic behaviour over time
 - Risk is not always at a single moment
 - Delays, accumulation of backlog, system degradation, human fatigue are all time-based phenomena
- Use systems thinking to inform mitigation
 - Once you see the interplay, you can design interventions that take into account people, process and systems together (not just patch one)

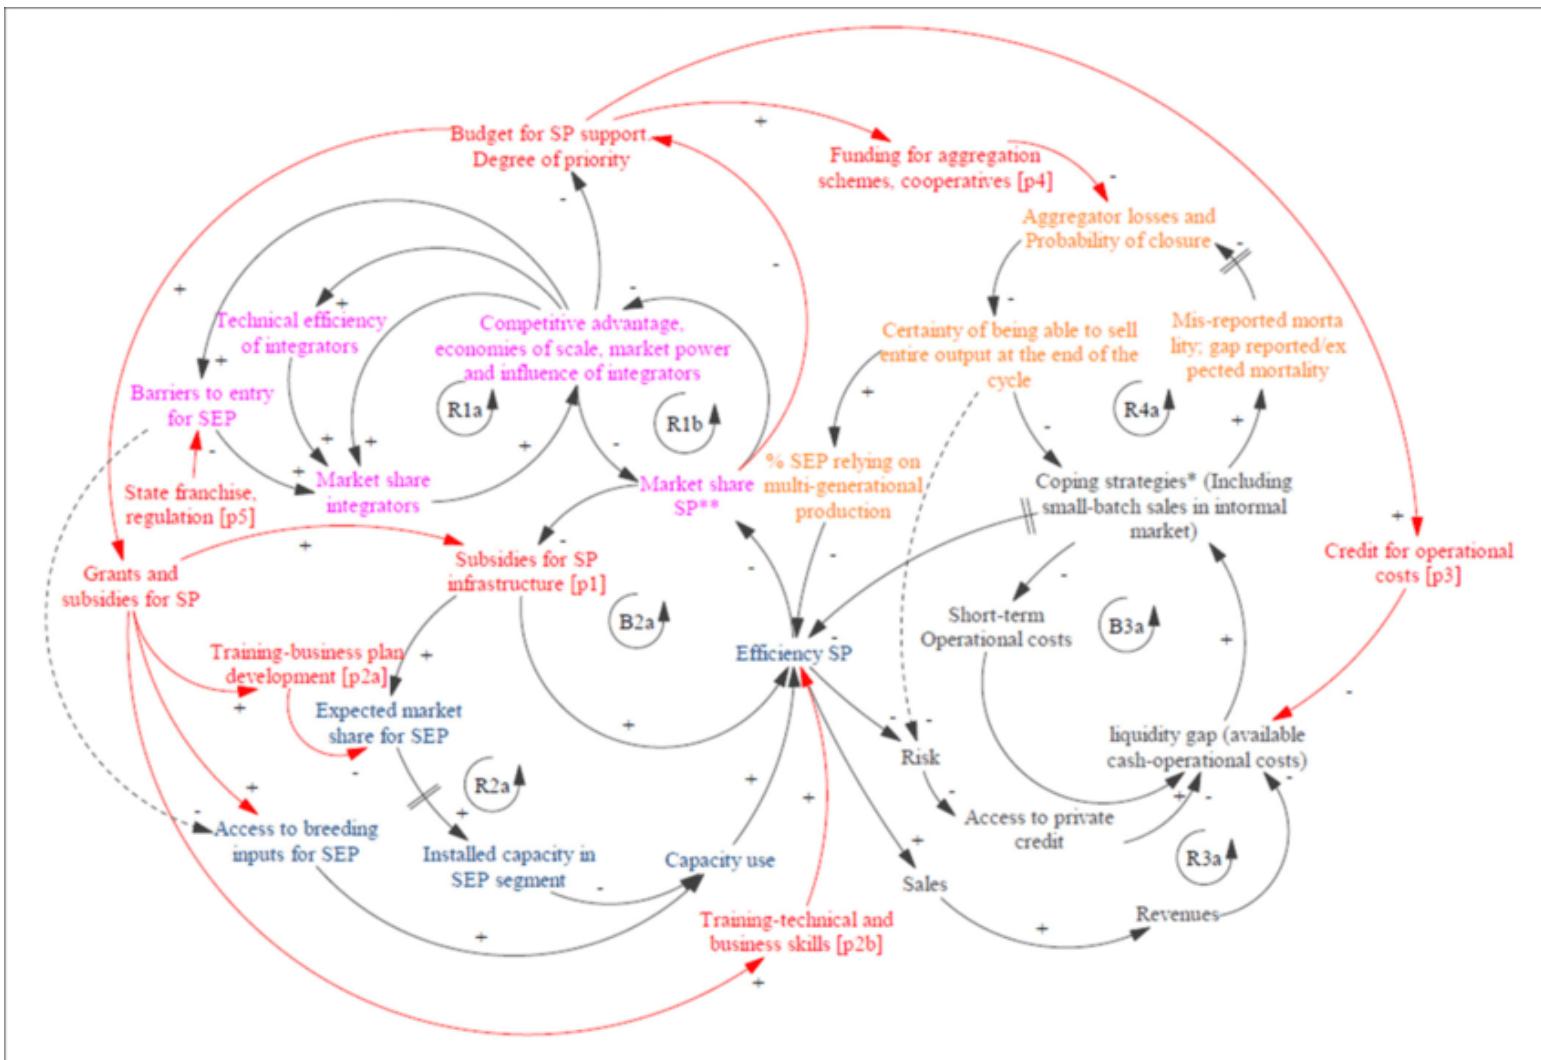
SYSTEMS THINKING TOOLS

- Causal loop diagram
 - Visualizes cause-and-effect chains and feedback loops
 - Useful to show how a process flaw or system delay may loop back and increase risk
- Behaviour over time graph
 - Chart of how a key variable changes over time
 - For example: backlog size, error rate, system performance

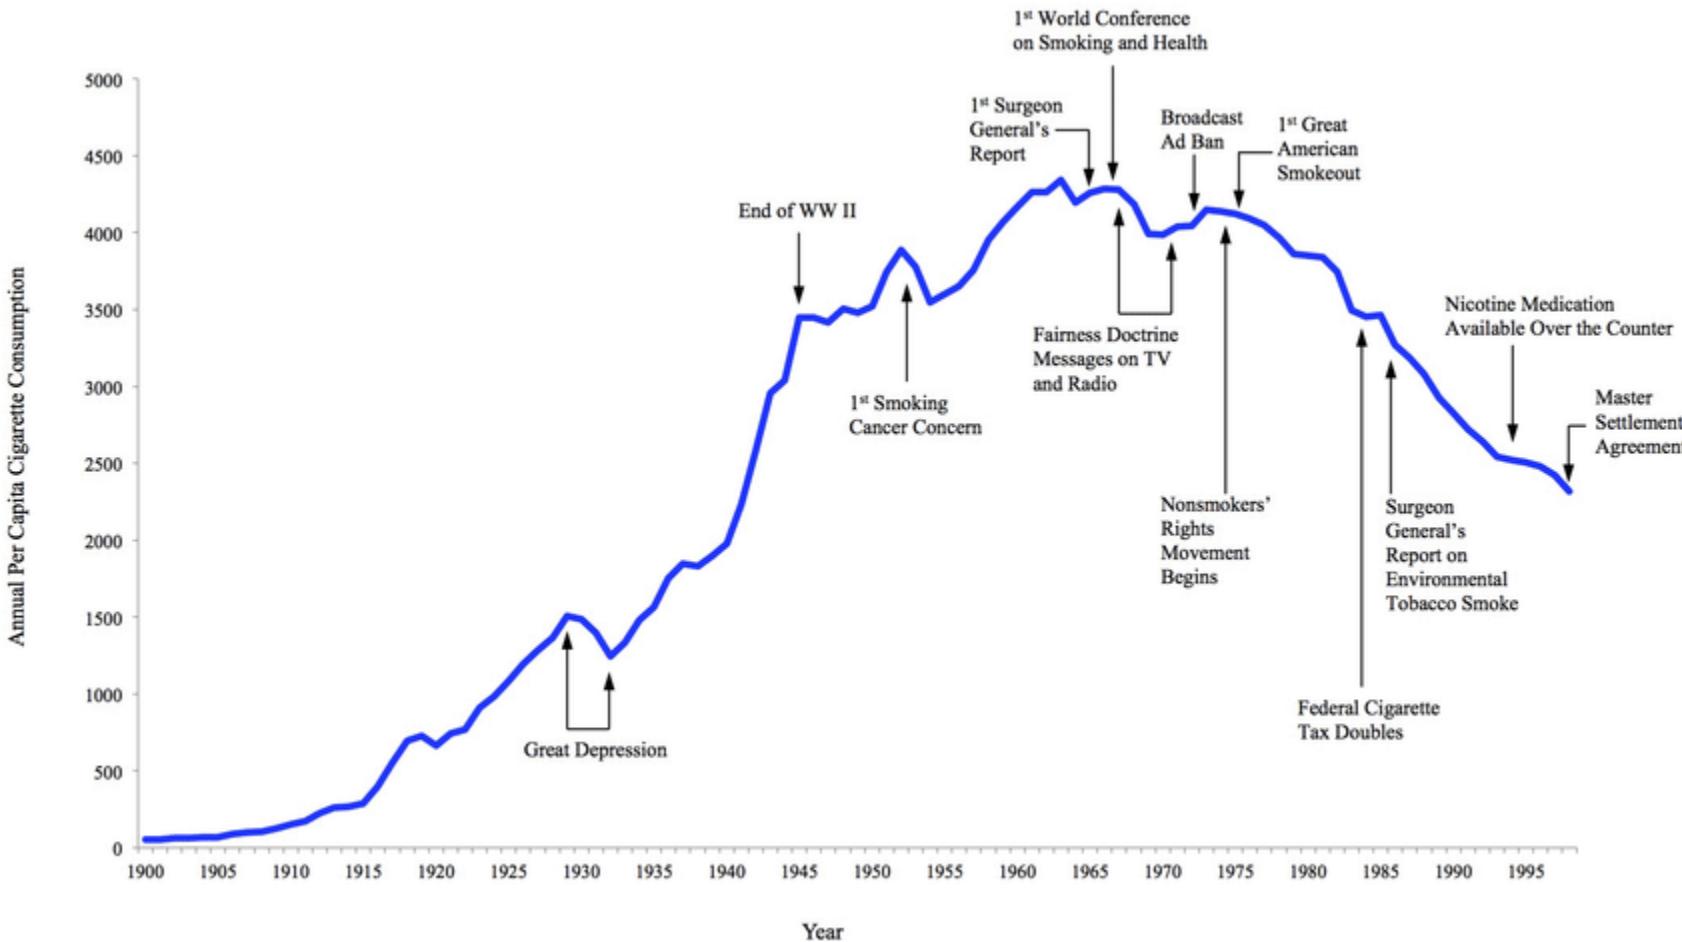
SYSTEMS THINKING TOOLS

- Stock and flow diagram
 - Shows accumulation (stock) and rates (flows) in processes
 - For example: tasks waiting in queues for processing
- Rich pictures
 - Qualitative depiction of system relationships, useful early in workshops to surface how people perceive the system
- Systems archetypes
 - Generic patterns like "Shifting the Burden", "Fixes that Fail", "Drifting Goals"
 - These help recognize recurring risk patterns in people-process-system interaction
 - Sort of like design patterns for identifying problematic systems architectures

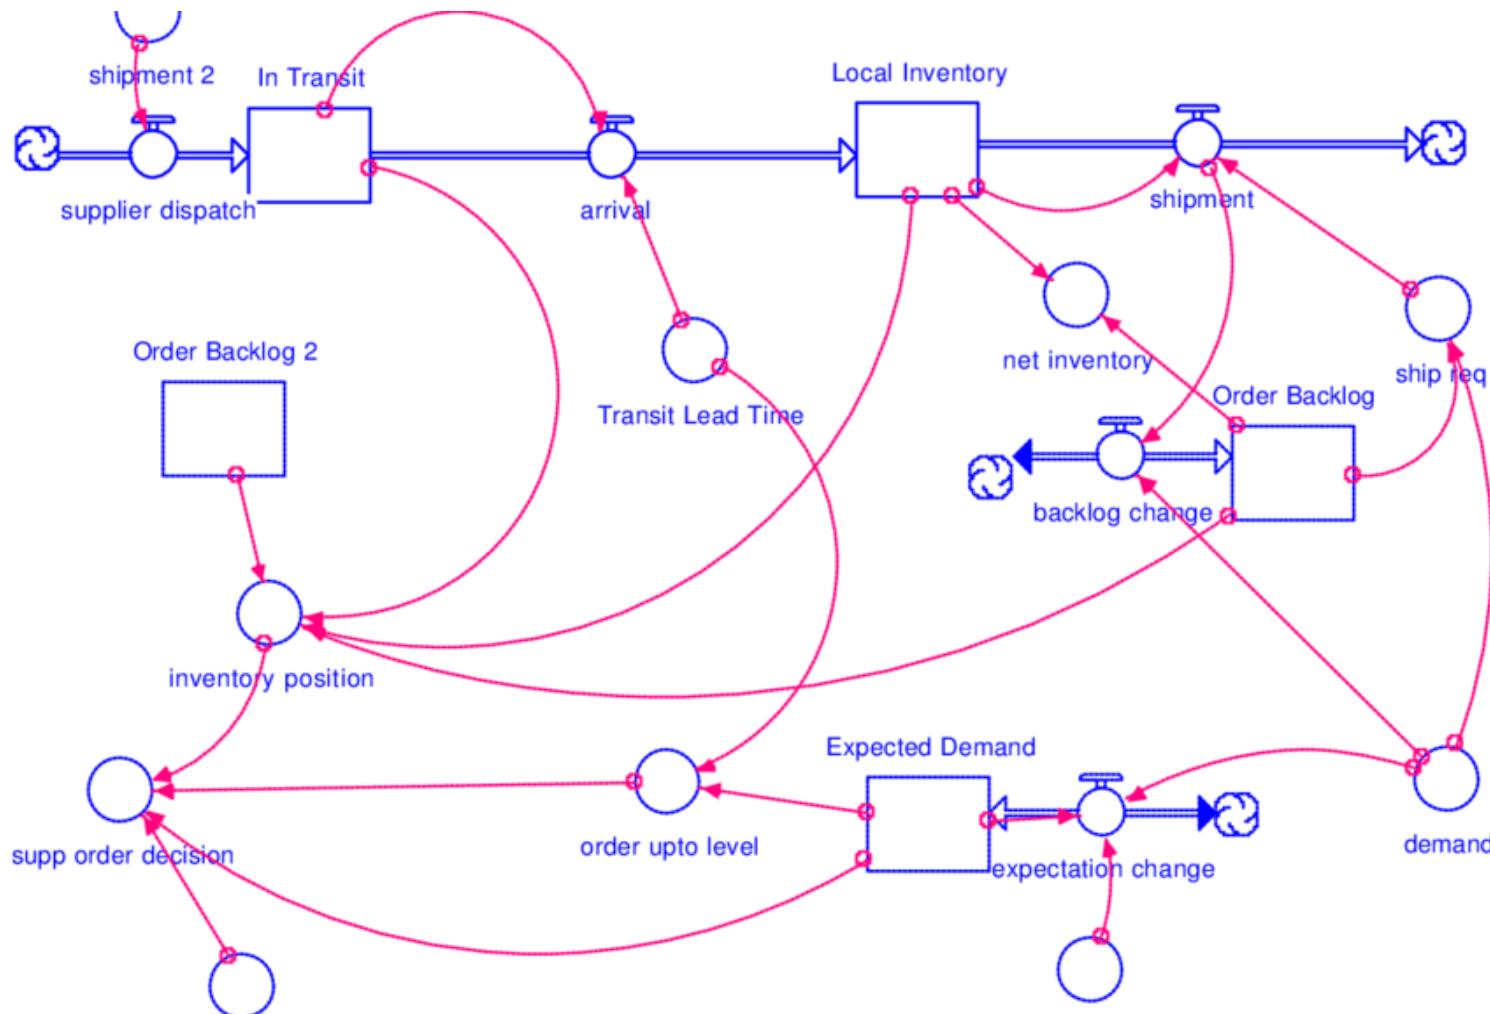
CAUSAL LOOP DIAGRAM



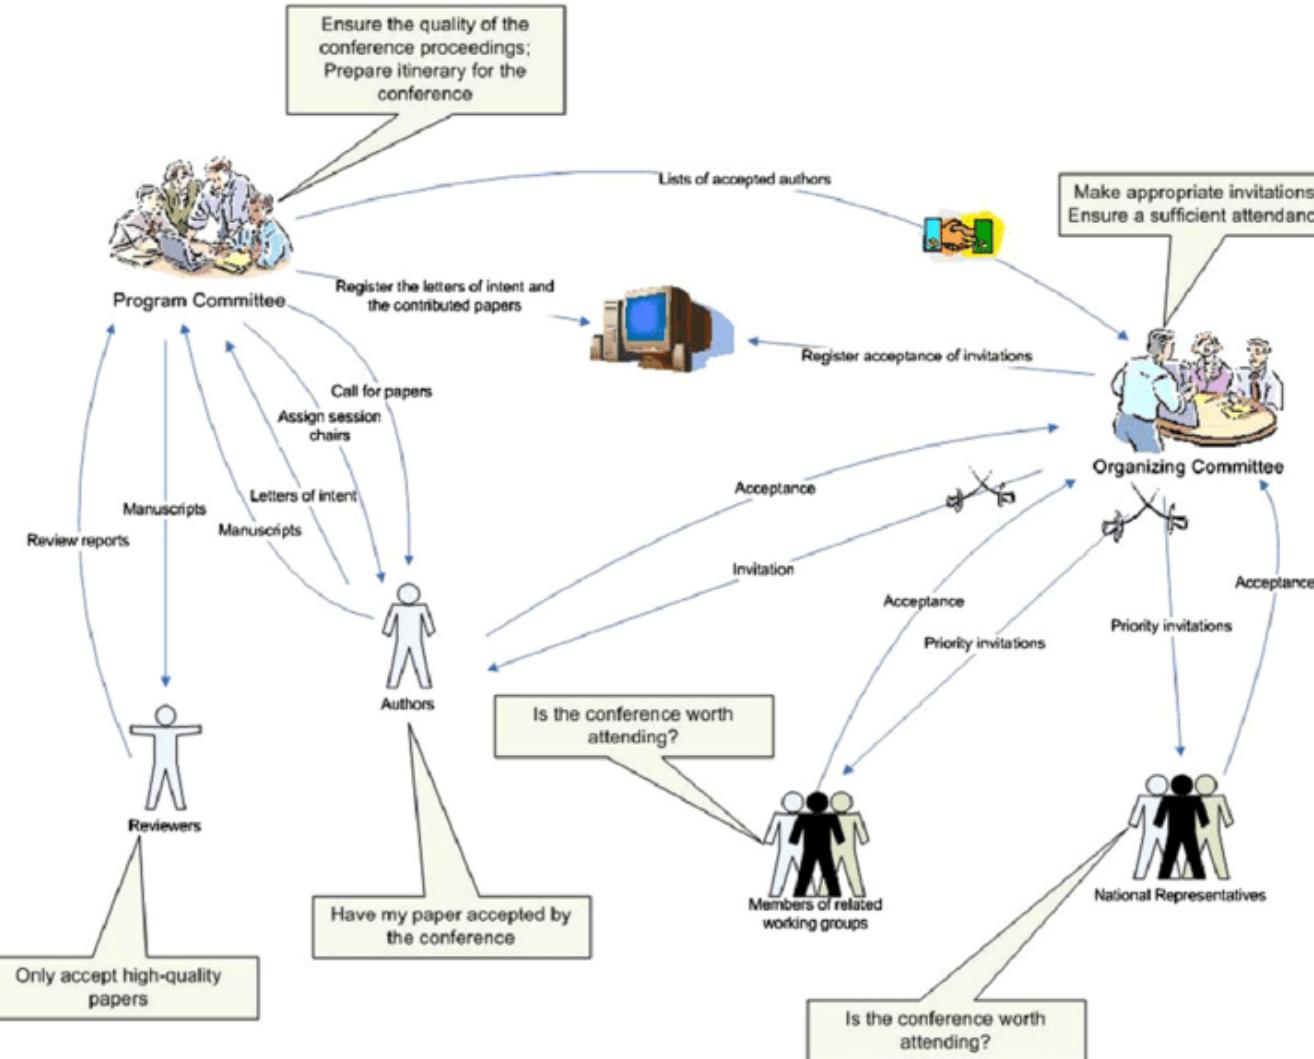
BEHAVIOUR OVER TIME GRAPH



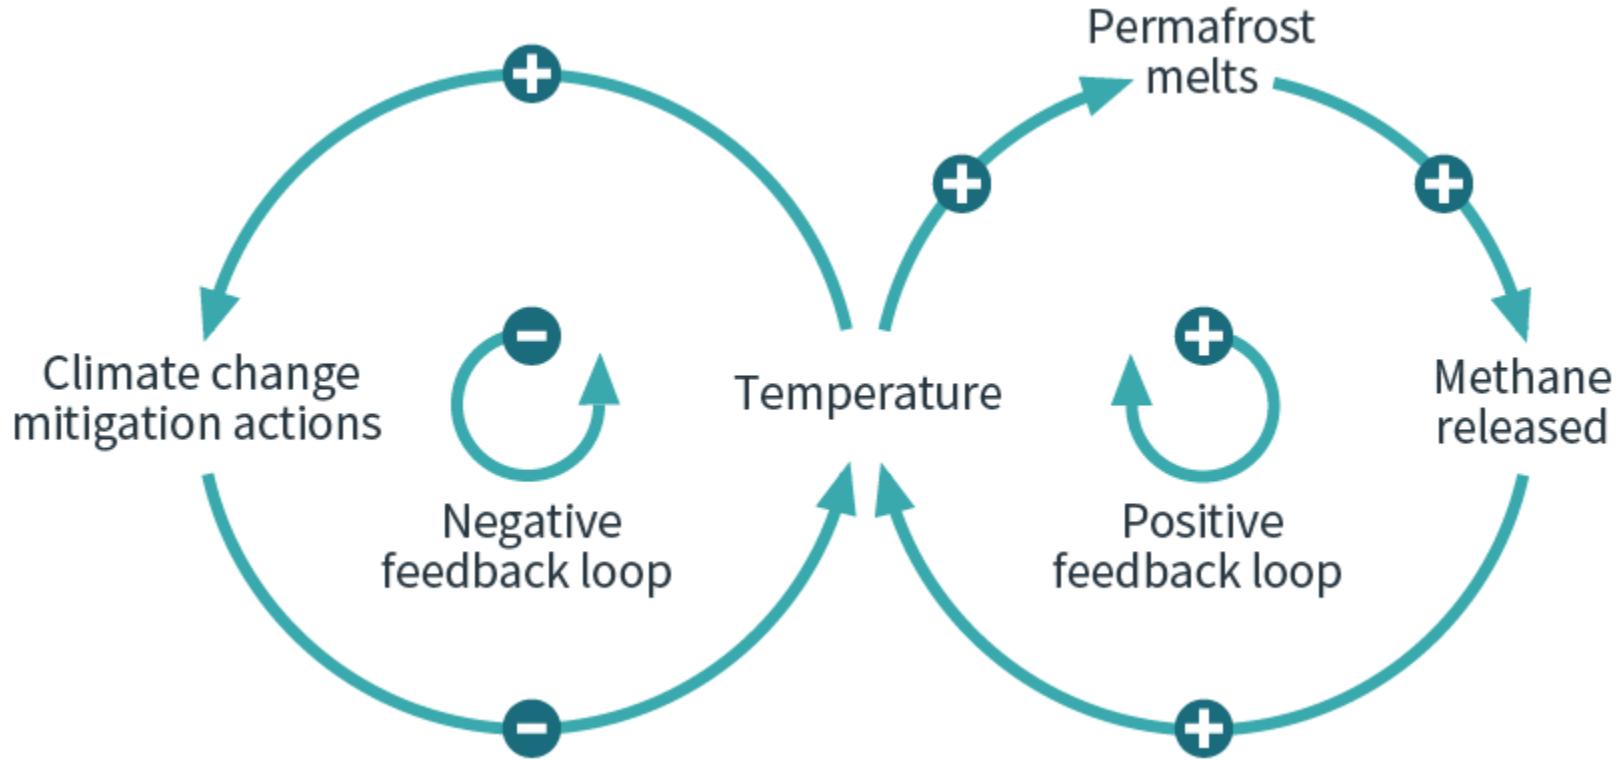
STOCK AND FLOW DIAGRAM



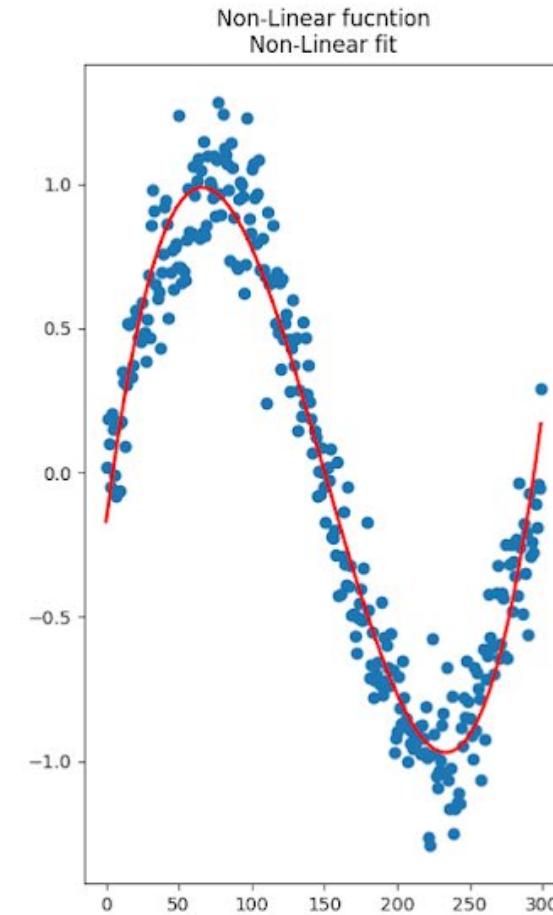
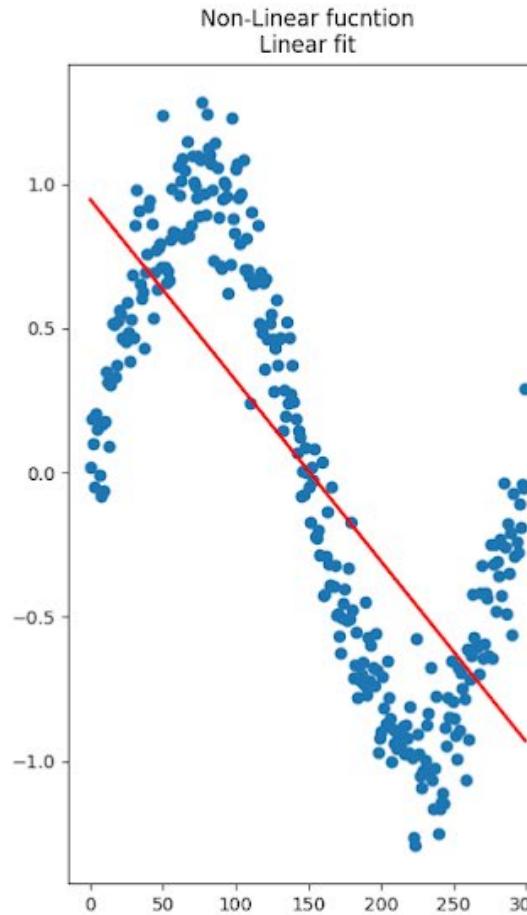
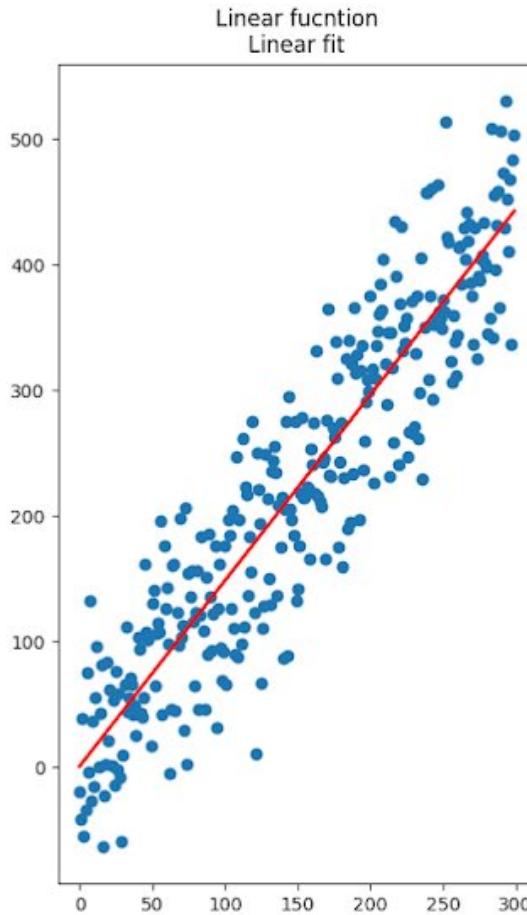
RICH PICTURE



FEEDBACK LOOPS



LINEAR VS NON-LINEAR



TIPPING POINTS

- A tipping point
 - The critical threshold at which a small change in one part of a system causes a sudden and often irreversible shift in the system's overall state or behaviour
- In complex systems
 - Feedback loops, inter-dependencies, and accumulated pressures can make a system appear stable right up until it reaches a tipping point
 - Beyond that threshold, the system rapidly reorganizes itself into a new pattern, often with disruptive consequences
 - Often seen in ecosystems, economies, and IT infrastructures
 - Basically a tipping point is where “more of the same” suddenly becomes “no longer the same”

TIPPING POINTS - KEY CHARACTERISTICS

- Non-linearity
 - Cause and effect are not proportional
 - Incremental stress builds quietly, but the response suddenly accelerates once the threshold is crossed
- Feedback-driven change
 - Reinforcing feedback loops amplify disturbances once they start
 - For example, error impacts workload performance that produces more error resulting in a loss of control
- Irreversibility or hysteresis
 - Returning conditions to their pre-tipping-point state doesn't always restore the system
 - New equilibria may form
 - For example: being locked into a crash recovery loop

TIPPING POINTS - KEY CHARACTERISTICS

- Sensitivity to small triggers
 - Minor perturbations like a delayed patch, missed communication can tip a system already under hidden stress into collapse
- Early-warning signals
 - Systems nearing tipping points often show growing volatility, slower recovery from small disturbances, or accumulating backlogs
 - Each warning signal on its own may not be enough to trigger a control
 - But when these signals together may be symptomatic of an emerging system instability

TIPPING POINTS - EXAMPLES

- Infrastructure overload
 - A system's performance declines gradually as load increases
 - But beyond a tipping point, latency spikes, processes hang, and cascading service outages occur
- Operational fatigue
 - Staff under constant high workload start making small mistakes
 - Error rates build until quality collapses
 - A human-process tipping point.
- Security risk escalation
 - Patch delays accumulate
 - Dependency vulnerabilities multiply
 - One exploit triggers a chain reaction that brings systems offline
- Cultural or governance shifts
 - Persistent blame culture leads to suppressed reporting until a major incident forces a governance overhaul

TIPPING POINTS - RELEVANCE TO RISK

- Monitoring leading indicators is as important as tracking lagging events
 - Traditional analysis assumes proportional change
 - Systems thinking reveals that risk often behaves non-linearly
 - Understanding tipping behaviour helps detect when an environment is approaching instability before catastrophic change occurs
- Early detection allows for preventive resilience
 - Adjusting load, redistributing responsibilities, or introducing dampening feedbacks before the threshold is crossed

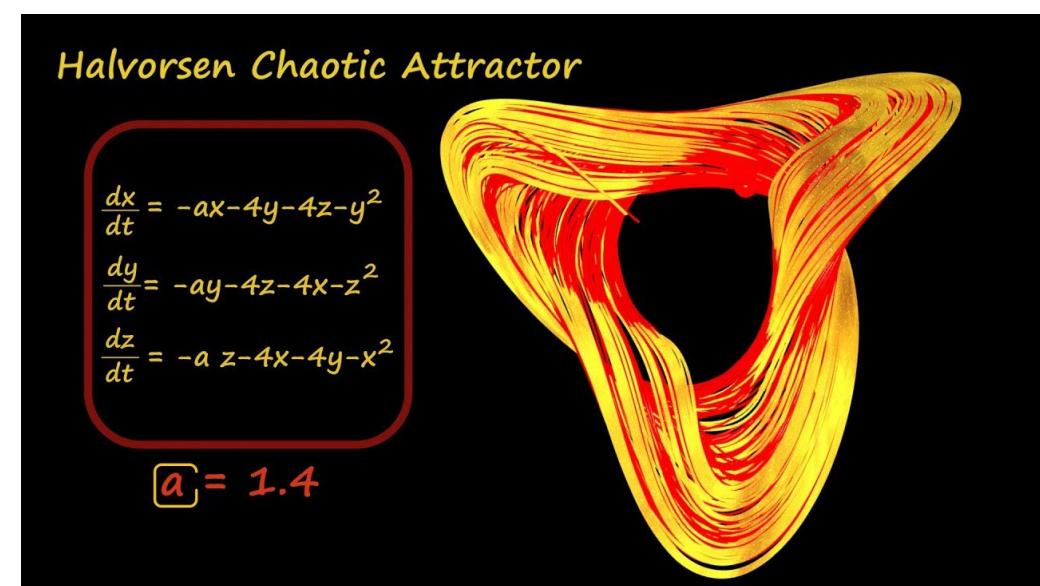
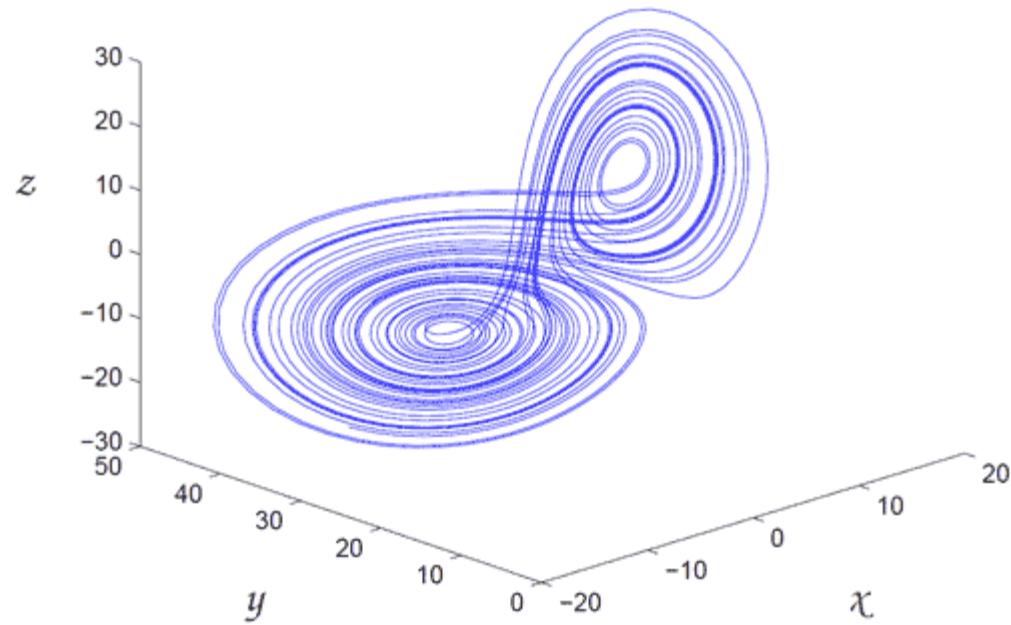
MITIGATION AND DETECTION STRATEGIES

- Identify critical thresholds
 - Use metrics (e.g., latency, backlog size, error rates, turnover) to define points where performance sharply degrades
- Monitor leading indicators
 - Look for signals such as increasing variability, repeated near-misses, or slower recovery from small disruptions
- Build buffers and redundancy
 - Introduce spare capacity or modular design so that stress is absorbed gradually rather than accumulating.
- Strengthen feedback controls
 - Implement mechanisms that balance reinforcing loops
 - For example: automatic throttling, escalation alerts, adaptive load balancing.
- Encourage transparent reporting
 - A culture that surfaces weak signals early prevents invisible accumulation of risk pressure

MITIGATION AND DETECTION STRATEGIES

- Tipping points
 - Are the dynamic consequence of feedback loops, delays, and interdependencies
 - When mapping causal loops, a tipping point is often where a reinforcing feedback loop overwhelms balancing controls
 - By identifying and modelling these loops, risk teams can anticipate thresholds and design interventions that restore balance before irreversible change occurs
- The formal study of tipping points is in the mathematical domain called chaos theory
 - Studies that some systems shift into configurations at tipping points
 - The configurations are called attractors
 - An IT system can go from stable to a number of unstable states
 - The effort required to revert back to a stable state can be very high

ATTRACTORS



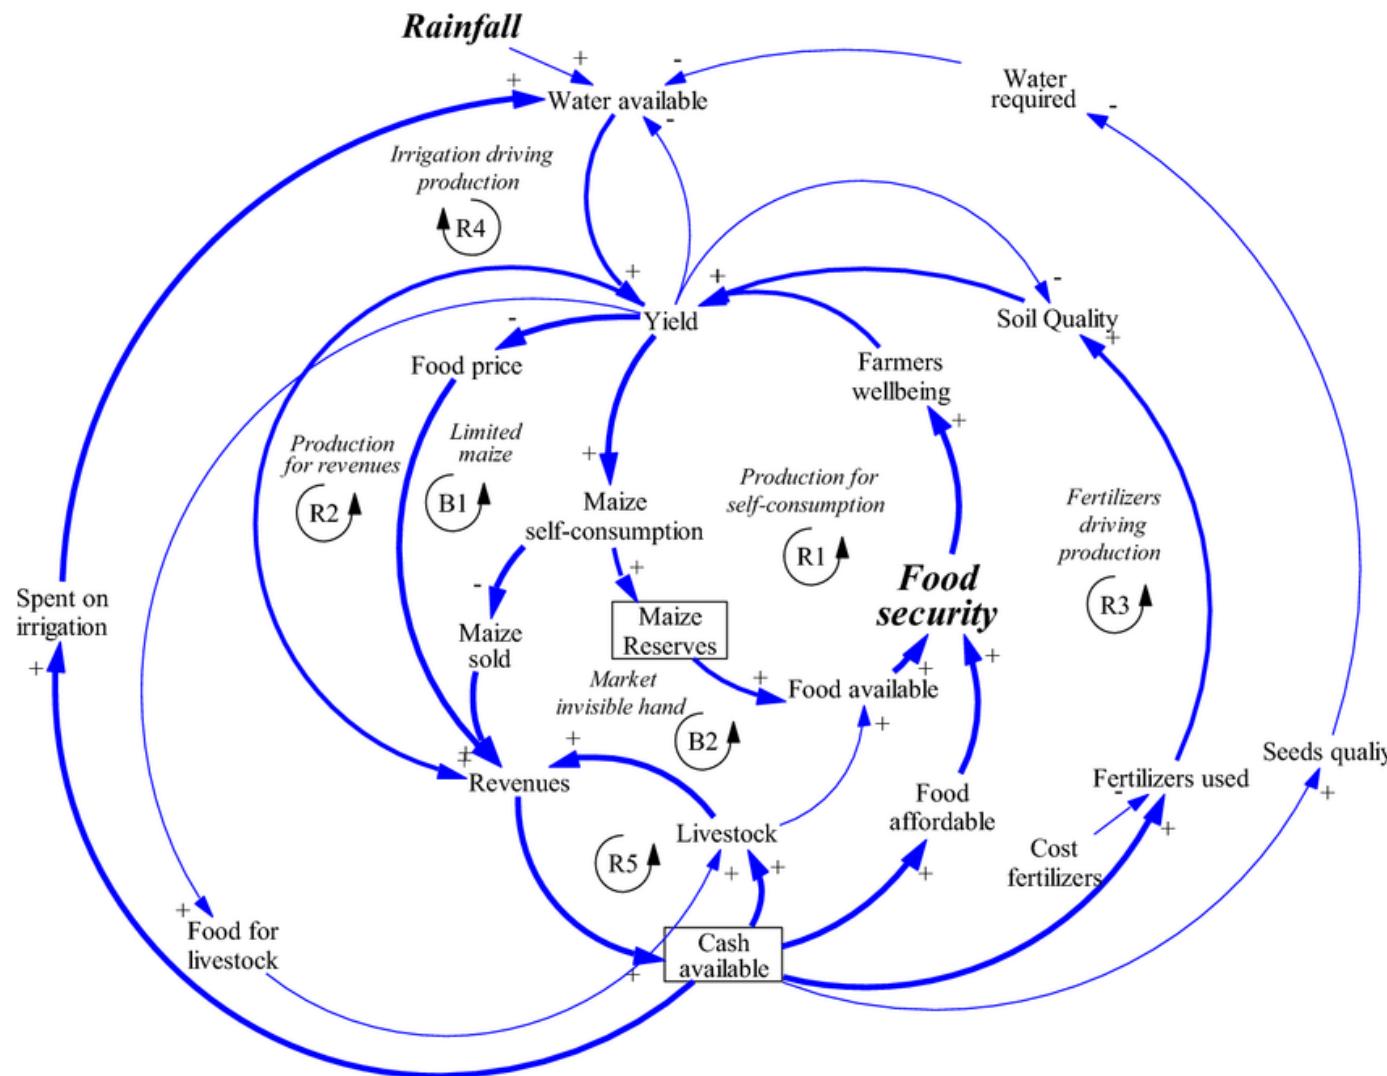
CASCADING FAILURES

- A cascading failure
 - Occurs when a failure or disruption in one component of a system triggers failures in other components
 - Often in sequence or via feedback loops, resulting in a large-scale systemic disruption.
 - A small error or delay in one process or team can propagate through downstream processes, across systems, or between teams
 - Effects are amplified by hidden dependencies or lack of redundancy
 - Example:
 - A key system goes offline
 - Manual workaround puts load onto a team
 - That team delays
 - Downstream system receives incorrect input
 - Compliance control is bypassed
 - Regulatory breach occurs.

HIDDEN DEPENDENCIES

- Dependencies are often not visible in standard process maps
 - For example, Team A's output becomes system input for Team B, but the dependency is informal/un-documented
- Hidden dependencies increase risk because
 - They create single points of failure (if Team A fails, Team B can't proceed)
 - They produce tight coupling (the downstream team or system has no buffer)
 - They reduce visibility and monitoring (failure upstream may not be flagged until it impacts downstream)
 - Systems thinking identifies these by asking
 - Which processes/systems/teams are reliant on another?
 - What would happen if that upstream function failed or was degraded?

CAUSAL LOOP DIAGRAM



CAUSAL LOOP DIAGRAM

- Purpose and overview
 - It helps visualize how different variables in a system interact and influence one another
 - Often through feedback loops that either reinforce (amplify) or balance (stabilize) change
 - CLDs shows how causes and effects circulate through a process
 - This reveals hidden dependencies, feedbacks, and potential tipping points
 - Instead of showing just a linear “A causes B,” a CLD shows networks of cause-and-effect relationships that evolve over time

KEY CONCEPTS

- Variables
 - Each node in the diagram represents a variable representing something that can increase or decrease over time
 - Examples
 - System Load
 - Team Workload
 - Error Rate
 - Backlog Size
 - Control Effectiveness
 - Each variable is connected by arrows that indicate causal influence

KEY CONCEPTS

- Causal links (arrows)
 - Arrows show direction of influence, how one variable affects another
 - Each arrow has a polarity, marked with a “+” or “–” sign
 - + (same direction)
 - When A increases, B increases
 - When A decreases, B decreases
 - For example: Workload (+) → Stress: means that workload increases stress
 - – (opposite direction)
 - When A increases, B decreases
 - When A decreases, B increases
 - For example: “Control Effectiveness (–) → Error Rate” means that better controls reduce errors.
 - Polarity does not mean good or bad, it just describes the direction of the effect

KEY CONCEPTS

- Feedback loops
 - A closed chain of causal links where a change in one variable eventually feeds back to influence itself
 - There are two major types
- Reinforcing (positive) feedback loop
 - Also called an R loop, this creates growth or amplification
 - Small changes compound over time, leading to exponential growth or runaway effects
 - Example:
 - More workload → More errors → More rework → More workload
 - Even if it starts small, the loop reinforces itself and grows until something breaks or stabilizes it or until it reaches a tipping point
 - Often labelled with R (or sometimes "+") in the loop symbol.
 - Or a small circle with the letter R at the loop center or end

KEY CONCEPTS

- Balancing (negative) feedback loop
 - Also called a B loop, this creates stability or regulation
 - Changes are counteracted over time, pushing the system back toward a target, goal, or equilibrium
 - Rather than compounding, effects dampen the original change
 - Example:
 - More workload → More prioritization → Fewer tasks accepted → Reduced workload
 - If a change occurs, the loop works to limit growth or decline, keeping the system within acceptable bounds
 - Often labelled with B (or sometimes “–”) in the loop symbol
 - Or a small circle with the letter B at the loop center or end
 - They often include delays, which can cause oscillation or overshoot if not well managed

KEY CONCEPTS

- Time delays
 - Real systems rarely react instantly
 - To show this, a delay is marked with a small double slash “//” across a causal arrow
 - For example: “Incident Occurs → // → Management Awareness”
 - This means it takes time for information or consequences to propagate through the system, a crucial source of instability and surprise

HOW TO READ A CLD

- To interpret a causal loop diagram
 - Start with one variable and follow the arrows
 - Track the direction (same “+” or opposite “–”) of each causal link
 - Continue following until you return to the starting variable, that’s a loop.
 - Count the number of negative (–) signs in the loop
 - Even number → Reinforcing Loop (R)
 - Odd number → Balancing Loop (B)
- Ask
 - Does the loop amplify change (R) or dampen it (B)?
 - Are there delays that could make reactions too slow or too strong?
 - What risks or tipping points might arise if this loop becomes dominant?

EXAMPLE

- Scenario: incident backlog loop
 - Workload (+) → Stress (+) → Errors (+) → Rework (+) → Workload
 - This is a Reinforcing Loop (R1)
 - More work leads to more errors, which increases rework, which further increases workload
- To balance it, you might add a control
 - Errors (+) → Management Oversight (-) → Errors
 - Creates a Balancing Loop (B1)
 - Increased oversight reduces errors, stabilizing the system
 - If the reinforcing loop grows faster than the balancing loop can react (especially if there's a delay in oversight)
 - Then a tipping point may occur where the process collapses under workload stress

Q&A AND OPEN DISCUSSION

